

Online Algorithms and the Ski Rental Problem

1. Introduction to Online Algorithms

Online algorithms deal with the problem of decision making under partial information, where the inputs arrive as a stream of data, and at each time step, decision must be made on the fly, without knowing what will happen in the future. Examples of scenarios where online algorithms apply includes finding the shortest path for partially observable graphs (Canadian Traveler Problem), hiring high-quality applicants from interviews where individual decisions must be made right after the interview (Secretary Problem), and determining whether to rent or buy a particular resource at each time point (Ski Rental Problem).

We are interested in algorithms that guarantee a certain performance, or produce solution that is within a certain bounded ratio away from the optimal solution even in the worst case. This motivates the definition of competitive ratio.

Definition 1: *The competitive ratio of an online algorithm A is the worst case over possible future input sequences μ of the ratio $A(\mu)/A^*(\mu)$, where $A(\mu)$ represents the cost of A on μ and $A^*(\mu)$ is the cost of the optimal offline algorithm.*

2. The Ski Rental Problem Overview

Consider the following problem in online decision making. You came to a resort for skiing, uncertain about how many days you will ski. For each day in an ambiguous time period, you have to make decisions between renting or buying skis. Let's say that you can rent the skis at a cost of \$1 per day, or buy them for \$B. Once you buy the skis, you can ski rent-free for as long as you decide to stay, but before then you incur the daily rent cost. You would like to determine the "optimum" day on which you should buy skis.

We can formulate this into an online problem thus:

- Let x_d denote the input sequence where you ski for d days, and then stop
- Let A_i denote the deterministic algorithm where you rent the skis for $i-1$ days before buying the skis on the i -th day

Determining the optimum solution can be viewed as a game between you (the consumer) and an adversary that we will call nature who's able to determine when you would stop skiing.

- Your action set consists of $\pi_c = \{A_i : i \in [0, +\infty) \cup \{+\infty\}\}$ where $i = +\infty$ means you will always rent.
- Nature's action set consists of $\pi_n = \{x_d : d \in [0, +\infty) \cup \{+\infty\}\}$.

You first choose an algorithm A_i , then nature will choose the worst possible x_d based on your A_i so as to maximize $A(x_d)/A^*(x_d)$.

Assume that nature sets the time you ski to be d days (x_d) and you choose to buy the skis on day i (A_i), then the cost of A_i for different i 's is as follows:

$$C(A) = \begin{cases} B + i - 1 & \text{if } d \geq i \\ d & \text{if } d < i \end{cases}$$

While the cost of the optimal offline algorithm (A^*), with knowledge of d , is

$$C(A^*) = \min \{B, d\}$$

Thus the optimal algorithm chooses either to rent the whole time or buy from the beginning.

We want to minimize the competitive ratio ρ :

$$\text{competitive ratio } \rho = \max_{\text{over } d} \frac{C(A)}{C(A^*)}$$

One strategy we could choose is the “better-late-than-never” strategy, where we rent until we realize we should've bought, then buy. This means renting the ski for B days, then buy. Since nature wants to maximize the competitive ratio, it will choose to end to skiing time right after you buy skis. So:

- If $i \leq B$, then $\rho = \frac{i + B - 1}{i}$
- If $i > B$, then $\rho = \frac{i + B - 1}{B}$

We need to pick i to minimize ρ . This minimum is achieved when $i = B$, which gives us a competitive ratio of $\rho = 2 - \frac{1}{B}$.

To generalize this result, assume renting the ski costs $\$r$ per day, and buying the ski costs $\$b$, then:

Theorem 2: *The algorithm better-late-than-never has the best possible competitive ratio of $2 - \frac{r}{p}$ out of all the deterministic algorithms for the ski-rental problem.*

3. Randomized Algorithms for Ski Rental

Now we consider the extension where our action space contains probabilistic strategies and we make decisions based on coin tosses. Now, the adversary nature only knows our strategy but not the exact output. Now a strategy in our action space can be buying skis on day i with $P(i)$, where $\sum_{i=0}^{\infty} P(i) = 1$. Then, if the time of our ski is d , the cost of the optimal offline solution is still just $C(P^*) = \min \{B, d\}$. Now for our strategy $P(i)$, if we buy skis on day i , then we will need to pay $\$(i - 1 + B)$ if $i < d$ and $\$d$ if $i \geq d$. So the cost for our online algorithm becomes:

$$C(P(i)) = d \sum_{i \geq d} P(i) + \sum_{i < d} (i + B - 1) P(i) \approx d \int_d^{\infty} P(i) di + \int_0^d (i + B - 1) P(i) di$$

Our goal is:

Minimize $\rho = C(P(i)) / C(P^*)$ such that

$$\int_0^{\infty} P(i) di = 1$$

We have that $d \int_d^{\infty} P(i) di + \int_0^d (i + B - 1) P(i) di = \rho \min\{B, d\}$

- Case $d \geq B$: Differentiate with respect to d on both sides (applying chain rule on the left), we have

$$\begin{aligned} -dP(d) + \int_d^{\infty} P(i) di + (d + B - 1)P(d) &= 0 \\ (B - 1)P(d) + \int_d^{\infty} P(i) di &= 0 \end{aligned}$$

We can assume that $B > 1$ since otherwise it would just make sense to always buy before day 1, which is trivial. Under this assumption, we have $B - 1 > 0$ and $P(i) > 0$ for values of $i \geq 0$. Thus,

- Case $d < B$: Case $d \geq B$: Differentiate with respect to d on both sides (applying chain rule on the left), we have

$$\begin{aligned} -dP(d) + \int_d^{\infty} P(i) di + (d + B - 1)P(d) &= \rho \\ (B - 1)P(d) + \int_d^{\infty} P(i) di &= \rho \end{aligned}$$

Differentiating again gives:

$$B \frac{dP(d)}{dd} - P(d) = 0$$

$$P(d) = C e^{\frac{d}{B}}, \text{ for some constant } C$$

To solve for C , we plug $P(d) = C e^{\frac{d}{B}}$ into $\int_0^{\infty} P(i) di = 1$:

$$\int_0^{\infty} P(i) di = \int_0^B C e^{\frac{d}{B}} dd = 1$$

$$C = \frac{1}{B(e - 1)}$$

Thus

$$\rho = \frac{e}{e - 1}$$

4. Multi-Shop Ski Rental Problem

Now we generalize the ski rental problem and consider the case when there are multiple shops offering rental and purchase options for skis at difference prices. Now, we (the customer) must make decisions for not only where, but also when to buy skis. We want to find the optimum online mixed strategy from the customer's perspective. For this section, we restrict our analysis to the multi-shop ski rental problem (MSR), where the customer must choose one shop immediately after arrival at the ski resort and cannot change the shop once she chooses one.

Assume there are n shops in total denoted by $[n] = \{1, 2, 3, \dots, n\}$. Each shop j offers rental at the price of r_j and purchase at the price of b_j . We assume that $0 < r_1 < \dots < r_n$ and $b_1 > \dots > b_n > 0$ to ensure that there is no one shop that's always the optimal solution.

Again, we model our problem as a zero-sum game between the customer and nature. The number of times the customer skis is determined by nature. We want to analyze the optimal strategy for the customer. We begin the analysis by showing that only one shop has a positive buying probability at each time t . Then we show that the positive probability over a certain shop lies in a continuous interval, such that we can partition the time space into different intervals each corresponding to a different optimal strategy. At last, we explain a linear time algorithm for finding these breakpoints and also consider some variations of the MSR problem.

4.1 Formulation

The action of the customer can be represented by the pair (j, x) , where j is the index of the shop chosen by the customer, and x is the time when she chooses to buy skis. Thus the action set of the customer is

$$\psi_c \triangleq \{(j, x): j \in [n], x \in (0, +\infty) \cup \{+\infty\}\}$$

Here, $x = +\infty$ indicates the strategy where the customer always rent the skis and never buys.

We denote by y the time when nature stops customer from skiing. Thus the action set of nature is

$$\psi_n \triangleq \{y: y \in (0, +\infty) \cup \{+\infty\}\}$$

A fixed strategy profile can be denoted by $\langle (j, x), y \rangle$ and the cost associated with this profile is

$$c_j(x, y) = \begin{cases} r_j y & \text{if } y < x \\ r_j x + b_j & \text{if } y \geq x \end{cases}$$

Now, the customer's mixed strategy is a probability density over the set of fixed strategy (j, x) pairs denoted by $p \triangleq (p_1, \dots, p_n)$ where p_j is the density assigned to the strategy (j, x) for all feasible (j, x) . Furthermore, by properties of probability density function, we require that $\sum_{j=1}^n \int_0^\infty p_j(x) dx = 1$. The probability densities of nature can be denoted by $q(y)$ where we require $\int_0^\infty q(y) dy = 1$.

For a given p, y pair, the expected cost is:

$$C(p, y) \triangleq \sum_{j=1}^n c_j(p_j, y)$$

$$c_j(p_j, y) = \int_0^\infty c_j(x, y) p_j(x) dx$$

Now, given a mixed strategy pair $\langle p, q \rangle$, the competitive ratio is defined as:

$$R(p, q) \triangleq \int_0^\infty \frac{C(p, y)}{OPT(y)} q(y) dy$$

where $OPT(y)$ is the optimum offline solution that either rents from the first store for the entirety of y times or buys from the n th store in the beginning:

$$OPT(y) = \begin{cases} r_1 y & \text{if } y \in (0, B] \\ b_n & \text{if } y > B \end{cases}$$

Here, B generalizes the ratio between the buying price to the renting price and is defined as b_n/r_1 . Like in the previous sections, B determines the effective action sets of the customer and nature.

To simplify the constraints before formulating the objective, note that although we've defined x and y over the range $(0, +\infty) \cup \{+\infty\}$, the action set can actually be reduced greatly.

Lemma 3: *For nature, any strategy $y \in [B, +\infty)$ is dominated by the strategy of never stopping the customer. For the customer, any strategy (j, x) in which $x \in (B, +\infty) \cup \{+\infty\}, \forall j \in [n]$ is dominated by the strategy of buying at B in the same shop.*

Now, we can formulate our objective to be:

$$\text{minimize } \max_{y>0} \left\{ \frac{c(p,y)}{OPT(y)} \right\} \text{ subject to } p \in P$$

This is equivalent to:

$$\text{minimize } \lambda \text{ s. t.}$$

$$\frac{c(p,y)}{r_1 y} \leq \lambda$$

$$\sum_{j=1}^n \int_0^B p_j(x) dx = 1$$

$$p_j(x) \geq 0 \forall x \in (0, B], \forall j \in [n]$$

Furthermore, the inequality bound $\frac{c(p,y)}{r_1 y} \leq \lambda$, can be replaced with $\frac{c(p,y)}{r_1 y} = \lambda$ as we will show in the section right after.

4.2 Optimal Strategy Analysis

Let's first discuss some intuitions about the solution for MSR. If the buying time x is very small, then it makes more sense for us to prefer shop n since b_n has the minimum buying price. On the other hand, if x is larger, then we might prefer shop 1 since it has the lowest rent. In the optimal strategy, the times $[0, B]$ could be partitioned into sub-intervals where each interval is characterized by a single shop having positive possibility to be chosen by the customer. To formalize these intuition:

Theorem 4: The optimal strategy p^* satisfies:

- (a) There exists a constant λ such that $y \in (0, B)$, $\frac{c(p^*, y)}{r_1 y} = \lambda$
- (b) There exist $n + 1$ breakpoints: d_1, \dots, d_{n+1} such that $B = d_1 \geq d_n \geq d_{n+1} = 0$, and $\forall j \in [n]$, we have

$$p_j^*(x) = \begin{cases} a_j e^{r_j x / b_j}, & \text{if } x \in (d_{j+1}, d_j) \\ 0, & \text{otherwise} \end{cases}$$

in which a_j satisfies that $a_j b_j e^{r_j d_j / b_j} = a_{j-1} b_{j-1} e^{r_{j-1} d_{j-1} / b_{j-1}} \forall j = 2, \dots, n$

Proof (Sketch):

- (a) We can show that $\forall y \in (0, B)$, there exists a constant λ such that $\frac{c(p^*, y)}{r_1 y} = \lambda$. If $\frac{c(p^*, y)}{r_1 y}$ is not constant with respect to y , then we can find valley in this graph, where the

maximum points can be found on either right, left, or both, each case there is a better strategy, which contradicts p^* being the optimal strategy.

(b) We make use of the lemma that in the optimal strategy p^* , there exists $n + 1$

breakpoints: d_1, \dots, d_{n+1} such that $B = d_1 \geq d_n \geq d_{n+1} = 0$, which partition $[0, B]$ into n sub-intervals, such that $\forall j \in 1, \dots, n$ and $\forall x \in (d_{j+1}, d_j)$, $p_j^*(x) > 0$ and $p_i^*(x) = 0$ for any $i \neq j$. Thus we see that the index of the shop increases as x decreases. Taking the double derivative of the equality we have in (a) and solving for $p_j^*(x)$, we get:

$$p_j^*(x) = \begin{cases} a_j e^{\frac{r_j x}{b_j}}, & \text{if } x \in (d_{j+1}, d_j) \\ 0, & \text{otherwise} \end{cases}$$

We can write any p that satisfies (b) as

$$\frac{C(p, y)}{r_1 y} = \alpha_1 \frac{b_1}{r_1} e^{\frac{r_1 y}{b_1}} \forall y \in (0, B]$$

Then minimizing λ is the same as minimizing α_1 , since the rest of the terms in the equation above are constants.

Now we can formalize our objective in a different way:

$$\begin{aligned} & \text{minimize } \alpha_1 \\ & \text{subject to } \sum_{j=1}^n \alpha_j \frac{b_j}{r_j} (e^{\frac{r_j d_{j-1}}{b_j}} - e^{\frac{r_j d_j}{b_j}}) = 1 \\ & \alpha_j b_j e^{r_j d_j / b_j} = \alpha_{j-1} b_{j-1} e^{r_{j-1} d_{j-1} / b_{j-1}} \forall j = 2, \dots, n \\ & \alpha_j > 0 \forall j = [n] \\ & B = d_1 \geq d_n \geq d_{n+1} = 0 \end{aligned}$$

We can show that this problem reduces to solving for the endpoints (d_1, \dots, d_{n+1}) , since once we have the endpoints, we can derive a constant of proportionality Ω_j between α_j and α_1 so that $\Omega_j \alpha_1 = \int_{d_{j+1}}^{d_j} p_j(x) dx$. And then we can derive $\Omega = \sum_{j=1}^n \Omega_j$ s.t. $\Omega \alpha_1 = \sum_{j=1}^n \int_0^B p_j(x) dx = 1$. Then we can just solve for α_1 , and derive the entire solution from there. Now we just need to find a way to compute (d_1, \dots, d_{n+1}) .

4.3 Computing Breakpoints

We start by fixing $\alpha_1 = 1$, ignore the total probability sums up to 1 constraint, and try to find the (d_1, \dots, d_{n+1}) that maximizes $\sum_{j=1}^n \int_0^B p_j(x)dx$, which is just Ω since $\alpha_1 = 1$. This is equivalent to minimizing λ . Then we just normalize all the α_j by dividing them by Ω .

Let P_j be the probability sum from shop j to n :

$$P_j \triangleq \sum_{\tau=j}^n \left(\int_0^{d_j} p_{\tau}(x)dx \right)$$

Then $P_1 = \Omega$. To compute d_j , knowing that d_1, d_2, \dots, d_{j-1} are fixed, we just need to maximize over possible values of P_{j-1} . $P_{j-1}(d_j)$ is a strictly concave function, so when we want to choose the maximal point, we just keep searching until we reach shop k , such that $d_n = d_{n-1} = \dots = d_{k+1} < d_k$. Then d_n should be maximal:

$$d_n = \frac{b_n}{r_n} \ln \left(\frac{b_k r_n - b_n r_k}{b_n (r_n - r_k)} \right)$$

d_n is unrelated to d_{n-1} if $d_n < d_{n-1}$. After $d_n, d_{n-1}, \dots, d_{j+1}$, we show how to get d_j . If $j = n$, then we just use the equation above with $k = n - 1$. If $j \neq n$, we again try to find $\max_{d_j} P_{j-1}(d_j)$.

4.4 MSR Algorithm

Theorem 5: *There is an n algorithm for computing the unique optimal strategy of the consumer. The time and space complexity of the algorithm are linear.*

Please see code implementation in the python file.

5. Work Cited:

1. Ai, Lingqing, Xian Wu, Lingxiao Huang, Longbo Huang, Pingzhong Tang, and Jian Li. "The Multi-shop Ski Rental Problem." *SAO/NASA ADS: ADS Home Page*. N.p., 01 Apr. 2014. Web. 11 May 2017.
2. "Online Algorithms." *An Introduction to the Analysis of Algorithms* (2012): 101-22. Web.
3. "Competitive Analysis." *Fashion Forecasting* (n.d.): n. pag. Web.
4. N. *Online Algorithms: Ski Rental* (n.d.): n. pag. Web.

5. Ali Khanafer , Murali Kodialam , and Krishna P. N. Puttaswamy Coordinated Science Laboratory, University of Illinois at Urbana-Champaign, USA. “The Constrained Ski-Rental Problem and its Application to Online Cloud Cost Optimization”